

Limitations of Analyzing Self-Reported Grad Cafe Data

Pammi (JHED ID gemefie1) - Software Concepts, Module 3

The 50,000 rows in the applicants table are not a sample of graduate applicants; they are a sample of people who chose to post on Grad Cafe, which is a very different population. Nobody is required to submit a result, so the database only contains applicants who know the site exists, who are engaged enough with the admissions process to report an outcome, and who feel their result is worth sharing. That last point matters most. Posting an acceptance to a well-known program is socially rewarding, and posting a rejection from one is at least cathartic, whereas an unremarkable admit to a mid-sized regional program is rarely announced anywhere. My Question 11 result shows this directly: the ten universities with the most Fall 2026 entries are Stanford, Berkeley, Yale, Washington, Wisconsin, Princeton, Michigan, Chicago, MIT and Toronto - elite research universities that receive a small share of all graduate applications in the United States but dominate the site. Whole fields are skewed the same way: PhD applicants account for 33,156 of the entries against 14,680 for master's degrees, even though far more people enroll in master's programs than doctoral programs nationally, because the long, anxious PhD admissions cycle is exactly what drives people to refresh forums. The self-reported metrics amplify the selection. Only about 59 percent of entries report a GPA, only about 8 percent report any GRE component, and among those who do, the average quantitative score is 165.75 (Question 3) compared with an ETS test-taker mean around 157. It is far more plausible that people with strong scores are the ones who type them in - and that applicants to quantitative PhD programs at the schools above are over-represented - than that the true applicant population averages 166. The database therefore tells me a great deal about what a self-selected, high-achieving, elite-program-focused slice of applicants experienced, and very little about the typical graduate applicant, who is largely absent from it.

Even within that slice, the values are whatever an anonymous person chose to type, with no verification and no consistent schema, so the numbers must be read as approximate descriptions of the submissions rather than measurements of reality. The clearest example came out of Question 3: Grad Cafe labels its score badge simply "GRE", and about 60 percent of the applicants who filled it in entered a combined total such as 328 rather than a quantitative score, so a naive $\text{AVG}(\text{gre})$ returned 261.49. I had to restrict each average to values that are possible on the metric's scale, and even then placeholders like a GPA of 9.99 or an Analytical Writing score of 99.99 had to be excluded, which means the "provided" population for each metric is a further, non-random subset. Program names are free text - "Johns Hopkins", "JHU" and "John Hopkins" all appear - so counts like Question 7's 20 JHU

Computer Science master's entries depend on how generous the string matching is, and the LLM standardization used for Question 9 can only be as good as the model and the canonical list behind it. Term, nationality and status are also self-classified: 1,224 rows carry no nationality at all and are silently dropped from Question 2's denominator, and a "Wait listed" post may or may not have been updated when the final decision arrived, so acceptance percentages such as Question 5's 40.74 percent for Fall 2025 measure the mix of outcomes people reported at the moment they posted, not a program's admit rate. Nothing prevents duplicate posts by one applicant across several programs, joke entries, or deliberate misinformation, and the anonymity that encourages people to share is the same thing that makes every field unverifiable. What the database supports, then, are relative and internal statements - Question 10's finding that reported master's outcomes are accepted far more often (67.59 percent) than reported PhD outcomes (24.38 percent) is probably directionally right, because it compares two groups subject to similar posting incentives - while absolute claims about "average GPA of admitted students" or "acceptance rate at Stanford" cannot be made from it. The correct reading of any figure here is "among people who posted this on Grad Cafe", and the responsible use of the data is for exploring patterns and generating hypotheses that would need a properly sampled source to confirm.